

Lecture Notes - 04: Logistic Regression

Dihui Lai

March 18, 2020

Contents

1 Likelihood Function

If a set of random variables $Y_1, Y_2 \dots Y_n$ has a joint probability distribution density/mass $f(y_1, y_2, \dots y_n; \theta)$, where θ is a set of parameters, the likelihood function is defined as

$$L(\theta) = f(y_1, y_2, \dots y_n; \theta) \quad (1)$$

Assuming an event has two possible outcomes $y = 1$ or $y = 0$, with probability p of being 1, i.e. the outcome follows a Bernoulli distribution. As we learned in lecture 2, the probability mass function is

$$f(y; p) = \begin{cases} p, & y = 1 \\ 1 - p, & y = 0 \end{cases}$$

Or

$$f(y; p) = p^y (1 - p)^{1-y}$$

The probability mass distribution (or the likelihood function by definition) for n independent events is

$$L(p_1, p_2, \dots p_n) = f(y_1, y_2, \dots y_n; p_1, p_2, \dots, p_n) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

To interpreting the likelihood function, let us consider the underlying parameters are the same i.e. $p = p_1 = p_2 \dots = p_n$ for all the data entries observed. And we have the likelihood function as

$$L(p) = \prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$

Let us consider the following cases $n = 1$, $n = 2$ and any n . What kind of p that can maximize the likelihood function $L(p)$?

- $n = 1$ (1 observation): The likelihood function is $L(p) = p^y (1 - p)^{1-y}$.

Observations	$L(p)$	$L_{max}(p)$
$y = 0$	$L(p) = 1 - p$	$L_{max} = 1$ at $p = 0$
$y = 1$	$L(p) = p$	$L_{max} = 1$ at $p = 1$

- $n = 2$ (2 observations): The likelihood function is $L(p) = p^{y^1+y^2}(1-p)^{(1-y^1)+(1-y^2)}$. Given the

Observations	$L(p)$	$L_{max}(p)$
$y^1 = 0, y^2 = 0$	$L(p) = (1-p)^2$	$L_{max} = 1$ at $p = 0$
$y^1 = 1, y^2 = 1$	$L(p) = p^2$	$L_{max} = 1$ at $p = 1$
$y^1 = 0, y^2 = 1$	$L(p) = p(1-p)$	$L_{max} = 0.25$ at $p = 0.5$

- $n = n_1 + n_0$ (n observations with n_1 1s and n_0 0s): The likelihood function is $L(p) = p^{n_1}(1-p)^{n_0}$. The likelihood function is maximized when

$$\frac{\partial \ell}{\partial p} = 0, \text{ where } \ell = \log(L(p)) = n_1 \log(p) + n_0 \log(1-p) \quad (2)$$

Solve equation (3) for p , we have

$$\begin{aligned} \frac{\partial \ell}{\partial p} &= \frac{n_1}{p} - \frac{n_0}{1-p} = 0 \\ \Rightarrow n_1 - n_1 p - n_0 p &= 0 \\ \Rightarrow p &= \frac{n_1}{n_1 + n_0} \end{aligned}$$

Overall, p maximize the likelihood function when it takes the value of the mean of observed y s

2 Logistic Regression

2.1 Likelihood Function

In general, every events could have its own underlying parameter p . For n -independent events, let us assume the parameters are p_1, p_2, \dots, p_n respectively. The corresponding log-likelihood function is thus

$$\ell(p_1, p_2, \dots, p_n) = \sum_{i=1}^n (y^i \log(p_i) + (1 - y^i) \log(1 - p_i)) \quad (3)$$

The log-likelihood function is the defined as the log transformation of the likelihood function

$$\ell = \log(L) = \sum_{i=1}^n y^i \log(p_i) + (1 - y^i) \log(1 - p_i) \quad (4)$$

2.2 Parameter Model

The parameter p_i is modeled as a logistic function of a set of m predictors $x_1^i, x_2^i, \dots, x_m^i$ or \vec{x}^i in vector notation.

$$p_i = \frac{1}{1 + \exp(-\vec{\beta} \cdot \vec{x}^i)} \quad (5)$$

2.3 Maximum Likelihood Estimation

The optimal model chooses β s that maximize the likelihood function ℓ , at the optimal point β s satisfy the following equations.

$$\frac{\partial \ell}{\partial \beta_j} = 0, j = 1, 2, \dots, m \quad (6)$$

Use ℓ 's definition in equation (4) and formula (5), we have

$$\begin{aligned} \ell &= \sum_{i=1}^n y^i \log \frac{p_i}{1 - p_i} + \log(1 - p_i) \\ &= \sum_{i=1}^n y^i (\vec{\beta} \cdot \vec{x}^i) - \log(1 + \exp(\vec{\beta} \cdot \vec{x}^i)) \end{aligned}$$

Insert it into equation (6), we have

$$\frac{\partial \ell}{\partial \beta_j} = \sum_{i=1}^n \left(y^i - \frac{1}{1 + \exp(-\vec{\beta} \cdot \vec{x}^i)} \right) x_j^i = 0, j = 1, 2, 3, \dots, m$$

To get the optimal β s, we need to solve the equation set. However, it is hard to do analytically, because of the nonlinear terms that contain $\beta \frac{1}{1 + \exp(-\vec{\beta} \cdot \vec{x}^i)}$. However, we can solve the problem numerically, using Newton-Raphson method.

3 Newton-Raphson Method

3.1 Single Variable

Consider a log-likelihood function of one parameter $\ell(\beta)$. In general, ℓ can be of any function and complex. With the hope that its derivative ℓ' is simpler, we use Taylor expansion for approximation around some point β_0

$$\ell(\beta) \sim \ell(\beta_0) + \ell'(\beta_0)(\beta - \beta_0) + \frac{1}{2} \ell''(\beta_0)(\beta - \beta_0)^2 \quad (7)$$

The derivative of equation (7) is thus

$$\ell'(\beta) \sim 0 + \ell'(\beta_0) + \ell''(\beta_0)(\beta - \beta_0) \quad (8)$$

The β^* that minimizes the log-likelihood function have $\ell'(\beta) = 0$ at the point i.e. $\ell'(\beta)|_{\beta=\beta^*} = 0$. Using equation (8), we have

$$\ell'(\beta_0) + \ell''(\beta_0)(\beta^* - \beta_0) = 0 \quad (9)$$

$$\Rightarrow \beta^* = \beta_0 - \frac{\ell'(\beta_0)}{\ell''(\beta_0)} \quad (10)$$

Recall that this is only an approximation solution and β^* is not exactly the optimal point with an arbitrarily chosen β_0 . However, we can hope that equation (10) brings us a little closer to the optimal point. To get a more accurate solution, we will need to use equation (10) iteratively i.e.

$$\beta_{k+1} = \beta_k - \frac{\ell'(\beta_k)}{\ell''(\beta_k)}, \text{ until } |\beta_{k+1} - \beta_k| < \delta$$

Here, $|\beta_{k+1} - \beta_k| < \delta$ is the convergence condition and δ is tolerance level. δ is usually set as a small number. The algorithms says that we can stop the iteration if we are very close to the optimal point.

3.2 Multiple Variable

In the case where the log-likelihood function is dependent on multiple parameters $\ell(\beta)$, the Taylor expansion is

$$\ell(\beta) \sim \ell(\beta_0) + \nabla \ell(\beta_0)^T (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T \mathbf{H}(\beta_0) (\beta - \beta_0) \quad (11)$$

Here β is a $m \times 1$ column matrix and \mathbf{H} is the $m \times m$ Hessian matrix, defined as

$$\beta = \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_m \end{bmatrix}, \mathbf{H} = \begin{bmatrix} \frac{\partial^2 \ell}{\partial \beta_1^2} & \frac{\partial^2 \ell}{\partial \beta_1 \partial \beta_2} & \cdots & \frac{\partial^2 \ell}{\partial \beta_1 \partial \beta_m} \\ \frac{\partial^2 \ell}{\partial \beta_2 \partial \beta_1} & \frac{\partial^2 \ell}{\partial \beta_2^2} & \cdots & \frac{\partial^2 \ell}{\partial \beta_2 \partial \beta_m} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 \ell}{\partial \beta_m \partial \beta_1} & \frac{\partial^2 \ell}{\partial \beta_m \partial \beta_2} & \cdots & \frac{\partial^2 \ell}{\partial \beta_m^2} \end{bmatrix}$$

Apply the gradient against β on equation (11), we have

$$\nabla \ell = \nabla \ell(\beta_0) + \mathbf{H}(\beta - \beta_0) \text{ see Appendix.}$$

At the optimal point we want to have $\nabla \ell = 0$ i.e.

$$\begin{aligned} \nabla \ell(\beta_0) + \mathbf{H}(\beta - \beta_0) &= 0 \\ \Rightarrow \mathbf{H}^{-1} \nabla \ell(\beta_0) + (\beta - \beta_0) &= 0 \\ \Rightarrow \beta &= \beta_0 - \mathbf{H}^{-1} \nabla \ell(\beta_0) \end{aligned}$$

The Newton-Raphson algorithm for multivariate model is therefore

$$\beta_{k+1} = \beta_k - \mathbf{H}^{-1} \nabla \ell(\beta_k), \text{ until } |\beta_{k+1} - \beta_k| < \delta \quad (12)$$

4 Appendix

4.1 The Gradient of Equation (11)

Starting with equation

$$\ell(\beta) = \ell(\beta_0) + \nabla \ell(\beta_0)^T (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T \mathbf{H}(\beta_0) (\beta - \beta_0) \quad (13)$$

To simplify the equation, we introduce the notation $\Delta\beta = \beta - \beta_0$. It is easy to see the derivative of each element of $\Delta\beta$ against β_j has the following property

$$\frac{\partial}{\partial \beta_j} \Delta\beta_i = \delta_{ij} \quad (14)$$

Here, δ_{ij} is the Kronecker delta, having the property $\delta_{ij} = \begin{cases} 0 & \text{if } i \neq j, \\ 1 & \text{if } i = j. \end{cases}$

On the other hand, if we write the log likelihood using the elements in the matrices, we have

$$\ell(\beta) = \ell(\beta_0) + \sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} (\Delta \beta_a) + \sum_{a,b=1}^m \frac{1}{2} \Delta \beta_a H_{ab}(\beta_0) \Delta \beta_b \quad (15)$$

Let us look at each term on the R.H.S of the equation when we take the partial derivative of ℓ against β_j .

- The first term becomes 0 as it is constant $\nabla \ell(\beta_0) = 0$.
- In the second term, only $\Delta \beta_a$ is dependent on β and we have

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} \left(\sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} (\Delta \beta_a) \right) \\ &= \sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} \frac{\partial \Delta \beta_a}{\partial \beta_j} \\ &= \sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} \delta_{aj} \\ &= \frac{\partial \ell(\beta_0)}{\partial \beta_j} \end{aligned}$$

In matrix format, we have

$$\nabla \left(\sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} (\Delta \beta_a) \right) = \nabla \ell(\beta_0)$$

- the third term has two variables dependent on β $\Delta \beta_a$ and $\Delta \beta_b$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} \left(\sum_{a,b=1}^m \frac{1}{2} \Delta \beta_a H_{ab}(\beta_0) \Delta \beta_b \right) \\ &= \sum_{a,b=1}^m \frac{1}{2} \delta_{aj} H_{ab}(\beta_0) \Delta \beta_b + \sum_{a,b=1}^m \frac{1}{2} \Delta \beta_a H_{ab}(\beta_0) \delta_{bj} \\ &= \sum_{b=1}^m \frac{1}{2} H_{jb}(\beta_0) \Delta \beta_b + \sum_{a=1}^m \frac{1}{2} \Delta \beta_a H_{aj}(\beta_0) \\ &= \sum_{b=1}^m \frac{1}{2} H_{jb}(\beta_0) \Delta \beta_b + \sum_{a=1}^m \frac{1}{2} H_{ja}(\beta_0) \Delta \beta_a, \text{ use the fact that } H_{aj} = H_{ja} \\ &= \sum_{d=1}^m H_{jd}(\beta_0) \Delta \beta_d, \text{ (a, b are dummy indices, set them to be c)} \end{aligned}$$

In matrix format we have

$$\nabla \left(\sum_{a,b=1}^m \frac{1}{2} \Delta \beta_a H_{ab}(\beta_0) \Delta \beta_b \right) = \mathbf{H} \Delta \beta$$

Therefore we have

$$\nabla \ell(\beta) = \nabla \ell(\beta_0) + \nabla \left(\sum_{a=1}^m \frac{\partial \ell(\beta_0)}{\partial \beta_a} (\Delta \beta_a) \right) + \nabla \left(\sum_{a,b=1}^m \frac{1}{2} \Delta \beta_a H_{ab}(\beta_0) \Delta \beta_b \right) \quad (16)$$

$$= 0 + \nabla \ell(\beta_0) + \mathbf{H} \Delta \beta \quad (17)$$